

UNICODE

Stephen Schaub

Facts of Life

2

- FOL #1: Computers store and transmit data in units of bytes
- FOL #2: The world's languages require more than 256 characters
 - ▣ Single-byte encodings like ASCII map individual bytes to characters
 - ▣ All single-byte encodings pretend that FOL #2 doesn't exist

A Brief History of Encoding

3

- 1 byte per character schemes
 - ▣ ASCII
 - ▣ Code page systems
- 2 bytes per character schemes
 - ▣ Asian languages
- Unicode
 - ▣ 1-4 bytes per character

Unicode

4

- Maps characters to code points
 - ▣ A **code point** is a unique number that signifies a particular character
 - ▣ Current count: Over 137,000 code points / characters
- Code points 0-127 correspond to ASCII code

Unicode Planes

5

- Code points organized into 17 planes
 - ▣ Each plane can represent up to 65,535 code points
 - ▣ Plane 0: Basic multilingual plane (BMP)
 - Characters for almost all modern languages; several symbols
 - ▣ 16 supplementary "astral" planes
 - Historic languages
 - Music notation
 - Emoji
- Room for over 1 million code points
 - ▣ Most will likely never be assigned

Unicode Encodings

6

- Remember FOL #1? We need a way to represent code points using bytes.
- Various encodings possible
 - ▣ UTF-8
 - ▣ UTF-16
 - ▣ UTF-32
- UTF-32 can represent any possible code point in a single 4-byte value
 - ▣ Rarely used in practice (too inefficient)
- Both UTF-8 (1 byte values) and UTF-16 (2 byte values) are **variable-length** encoding systems
 - ▣ UTF-8 requires 1-4 bytes to represent a given code point
 - ▣ UTF-16 requires 2 or 4 bytes to represent a given code point

UTF-8

7

- Most widely used encoding of Unicode
- Requires 1-4 bytes to represent a given code point

48	69	e2	84	99	c6	b4	e2	98	82	e2	84	8c	c3	b8	e1	bc	a4
H	i	P			y		☂		§			ø		ñ			

HTML Page Encoding

8

- Content-type header can specify encoding
 - ▣ Content-Type: text/html; charset=utf-8
- Document charset meta tag can specify encoding

```
<html>
```

```
<head>
```

```
<meta charset="utf-8">
```

```
</head>
```

```
...
```


Practical Problems

9

- Variable-length encodings are memory-efficient, but not friendly for random access
 - ▣ Consider a string API that allows you to index a Unicode string:
 - `for (var i = 0; i < str.length; ++i)`
 `console.log(str[i])`
 - ▣ What are the options for implementing this behavior?

Practical Problems

10

- Some characters have multiple code point representations
 - ▣ Example: An accented e (é)
 - é (U+00E9)
or
 - e (U+0065) + accent (U+0301)
 - ▣ Example: The sequence fi
 - fi (U+FB01)
or
 - f (U+0066) + i (U+0069)
- Comparing Unicode strings for equality can be tricky

Case Folding

11

- Problem: Need to compare two strings to see if they contain the same letters, ignoring capitalization
 - ▣ With simple ASCII, case-insensitive comparisons are straightforward
 - ▣ Convert all letters in a string to the same capitalization (“case fold”) and then compare for equality

Case Folding

12

- With Unicode, the problem of comparing strings is more complex
- Three distinct issues:
 - ▣ Ignoring different Unicode representations for the same characters
 - Example: Compare “é” (U+00E9) to “é” (U+0065 U+0301)
 - ▣ Ignoring capitalization
 - Example: Compare “MASSE” to “Maße”
 - ▣ Ignoring accents, diacriticals, and other symbols
 - Example: Compare “Saens” to “Saëns”
- See https://www.w3.org/International/wiki/Case_folding

Handling Different Unicode Representations

13

- Unicode defines normalization forms and algorithms that convert two strings with different representations to the same representation

- Further reading:
 - ▣ See https://en.wikipedia.org/wiki/Unicode_equivalence

Handling Capitalization and Diacritical Differences

14

- Unicode defines algorithms for caseless matching
 - ▣ Libraries are available for various languages that implement this algorithm
- Other algorithms are available for removing accents

15

Case Studies

Searching Unicode Text

16

- BJU Digital Music Project
- Database contains composition titles and composer names like
 - ▣ Camille Saint-Saëns
- Problem: Searches need to match Saens == Saëns
- Solution:
 - ▣ Remove diacritical marks
 - ▣ Compare using case folding

Printing Source Code

17

- Code Listings Utility
- Problem: Print source code in unknown encodings
- Solution: Node detconv
 - ▣ Uses port of <https://github.com/chardet/chardet> to detect encoding
 - ▣ Uses <https://github.com/ashtuchkin/iconv-lite> to convert to ascii

18

JavaScript and Unicode

JavaScript and Unicode

19

- JavaScript represents strings internally using UCS-2
 - ▣ UCS-2 is a limited version of UTF-16 that handles only Plane 0 (BMP)
- Strings containing only Unicode values in the BMP (U+0000 to U+FFFF) often work well
 - ▣ Each of these can be represented using a single 16-bit value

Accented Characters

20

- é can be stored using one or two code points:
 - ▣ `const s1 = "\u00e9"; // é`
 - ▣ `const s2 = "\u0065\u0301"; // é`
- Both strings represent a single character
- JavaScript's `length` property counts code points, not characters
 - ▣ `s1.length == 1`
 - ▣ `s2.length == 2`
- Guess what you get when you index `s2[0]`? (Browser demo)

More JavaScript

21

- Consider emoji's 😊
 - ▣ Not in the BMP
 - ▣ Require a double-length UTF-16 value
 - ▣ JavaScript's string API exposes the underlying 16-bit representation in undesirable ways

Accurate indexing and character counts

22

- Unfortunately no good solutions exist at present in native JavaScript
- Node.js: Punycode library can help
 - ▣ Punycode represents Unicode in ASCII:
 - ▣ <https://en.wikipedia.org/wiki/Punycode>
- See <https://dmitripavlutin.com/what-every-javascript-developer-should-know-about-unicode/> for more suggestions

Comparing Strings

23

- Use the `.normalize()` method to convert two strings to a canonical representation that can be compared
 - ▣ `const s1 = '\u00E9' // é`
 - ▣ `const s3 = 'e\u0301' // é`
 - ▣ `s1 !== s3`
 - ▣ `s1.normalize() === s2.normalize()`
- Note that normalization preserves capitalization
 - ▣ The comparison is case sensitive
- No built-in way to do reliable case insensitive comparisons in JavaScript exists in 2020

Converting Bytes to Unicode

24

- Node.js: Byte data from files / sockets arrives as Buffer
- If data is textual, must convert Buffer to string
 - ▣ Specify encoding
 - ▣ `let str = buf.toString('utf8')`
- How do you know encoding?
 - ▣ Must be told
 - ▣ In general, not possible to infer

You Must Be Told the Encoding

25

- Cannot infer from a string of bytes; can only guess

	48	69	e2	84	99	c6	b4	e2	98	82	e2	84	8c	c3	b8	e1	bc	a4
utf-8	H	i	P			y		☂			§			ø		ñ		
iso8859-1	H	i	â	„	™	Æ	'	â	~	,	â	„	Œ	Ã	,	á	¼	¤
utf-16-le	橘		萌		呉		?		茈		萌		셀		?		X	
utf-16-be	輶		?		駟		등		颂		?		賃		룡		벤	
shift-jis	H	i	邃		卮		亓	筈		や		ゐ		て	ク	眈		,

26

Python and Unicode

Python and Unicode

27

- Python 3 has two sequence types that can hold textual info
 - ▣ bytes (unencoded byte values)
 - ▣ str (UTF-8 encoded Unicode)

```
>>> my_string = "Hi \u2119\u01b4\u2602\u210c\xf8\u1f24"  
>>> type(my_string) <class 'str'>  
>>> my_bytes = b"Hello World"  
>>> type(my_bytes) <class 'bytes'>
```

Normalization and Case Folding

28

- Normalize:
 - ▣ `import unicodedata`
 - ▣ `if unicodedata.normalize('NFC', s1) == unicodedata.normalize('NFC', s2)`
- Case insensitive compare:
 - ▣ Basic strings: `s1.lower() == s2.lower()`
 - ▣ Better: `s1.casefold() == s2.casefold()`
- Combining the two:
 - ▣ See <https://stackoverflow.com/a/40551443>



Pain Relief

Tip #1: Unicode Sandwich

30

- Bytes on the outside, Unicode on the inside
- Encode/decode at the edges
 - ▣ Receive binary data
 - ▣ Decode immediately to Unicode
 - ▣ Process as Unicode text
 - ▣ Encode as late as possible
 - ▣ Send binary data

Tip #2: Know What You Have

31

- Have a Buffer containing textual data?
 - ▣ Convert to string, specifying encoding
- Have a string and need a buffer?
 - ▣ Convert to buffer, specifying encoding

Tip #3: Avoid guessing the encoding

32

- Penalty for guessing wrong:
 - ▣ Some characters may fail to decode
 - ▣ The string may decode successfully, but produce wrong characters

Pro Tip #4: Test with unicode

33

- accéntéd téxt fǿř tēştīng
- Readable bυ ☂ η ☺ τ Α\$©!!
- iooł [n]θsn sı uMop-əpısdn

References

34

- Ned Batchelder, "Pragmatic Unicode."
<https://nedbatchelder.com/text/unipain.html>
- "What Every Programmer Absolutely, Positively Needs To Know About Encodings And Character Sets To Work With Text"
<http://kunststube.net/encoding/>
- Python 3 Unicode
<https://docs.python.org/3/howto/unicode.html>